

Optimization of Eigenface Algorithm for Facial Recognition Using AVX2

A Discussion Paper Presented to
the Faculty of the College of Computer Studies
De La Salle University

In Partial Fulfillment
of the Requirements for the Degree of
Bachelor of Science in Computer Science

by
Ibaoc, Christian Gabriel P.
Sia, Chantal Hyacynth S.
Valero, Nigel Kristoffer C.
Velooso, Antonio Luiz G.

Submitted to:
Professor Roger Luis Uy

July 30, 2024

Table of Contents

- I. Project Introduction
- II. Difference with Existing Implementation
 - A. asmNorm
 - B. asmNormDiv
 - C. subtnorm
 - D. covMat
 - E. jacobify
 - F. asmweights
 - G. eigPart1 and eigPart2
- III. Results and Conclusion

I. Project Introduction

In this project, the group implemented and optimized the Eigenface algorithm, an old but popular face recognition algorithm. Through provided images of a person's face, the algorithm can detect from its set of training images which image resembles the provided image the most. The base sequential implementation is done in C++, and the optimization is implemented in AVX2 assembly. Currently, existing parallelism implementations of the Eigenface algorithm is implemented in CUDA.

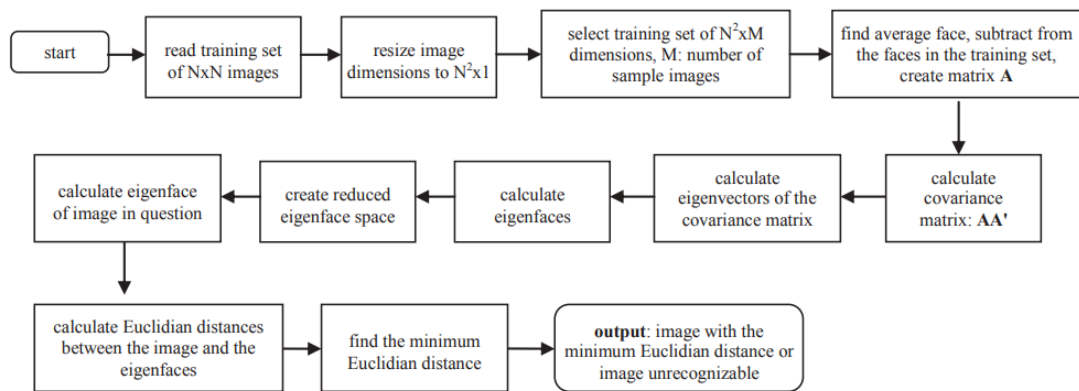


Figure 1.1. Flowchart of the Eigenface Algorithm [1]

The **main point** of the project was to produce some level of speed up to the base C++ algorithm by implementing the algorithm in AVX2 assembly. The implementation is scoped mostly to functions that perform some sort of matrix or vector calculation. The result is a level of speed up that is discussed in the following sections.

II. Difference with Existing Implementation

The algorithm has been implemented using CUDA [2] and MATLAB [3] before, so the main difference between this project's implementation is the usage of AVX2 and its instructions to optimize the algorithm and make it employ parallelism. The group also implemented the algorithm in C++ to compare the performance of their parallel implementation with a sequential one.

The `eigen.asm` file contains the implementation of the parallel Eigenface algorithm using AVX2, generally through vector operations on packed/ymm registers containing matrix pointers to be able to operate on more than one matrix element at a time. This differs from the sequential implementation, as it uses built-in functions to perform the steps of the algorithm and generally operates on one element at a time.

The following functions and explanations are used to implement parallelism to the Eigenface algorithm or to make specific steps of the algorithm more efficient:

A. `asmNorm`

`asmNorm` takes place in the “find average face” part of the algorithm, and it basically adds the pixel values of the training set array using ymm registers and packed vector addition until there is only one column left containing the sum of the training images, which will be divided to find the average image.

B. `asmNormDiv`

`asmNormDiv` takes place after `asmNorm` in the training phase, and it outputs the “average face” by dividing the sum of training images by the number of training images. Despite being a simple step at first glance, the AVX2 implementation is complicated by AVX2's limited ability to divide integers. This causes this function to require multiple steps of conversion (mostly to single-precision floating-point) to allow for packed vector division and to undo

previous conversions. Due to the complications caused by these conversions, the height and width of the training images are limited to integers divisible by 32.

C. subNorm

In this function, the “average face” given by `asmNormDiv` is simply subtracted from the values in the training set array. The packed vector subtraction doesn’t require prior conversions, so it is not as complicated as the previous function.

D. covMat

The covariance matrix is found by multiplying the normalized training set array with its transverse. This function takes in the normalized training set array and its transverse and performs matrix multiplication between them. By using vector packed multiply and add on packed registers containing the values of the training set array and its transverse, adding the results, and dividing it by the number of elements to process, the function outputs the value of a single normalized element of the covariance matrix at a row and column specified in the C++ file. As such, this function is called for each element of the covariance matrix.

E. jacobify

This function implements the Jacobian Rotation algorithm [4] to basically arrive at the eigenvalues and the eigenvectors from the covariance matrix. Specifically, the (estimated) diagonal of the output Jacobian matrix contains the eigenvalues and the matrix of eigenvectors (with column indices corresponding to the column index of the eigenvalues) is produced as an excess output.

Many of the steps for the Jacobian Rotation happen in the C++ file, but to keep the discussion related to AVX2 and parallelism, the `jacobify` function basically performs matrix multiplication on two given matrices. It does this through three levels of loops, the outermost of which (`greaterloop`) iterates over each row, the second of which (`jacobyloop`) iterates over each column, and the

innermost of which (innerloop) finds the value of a single element in the product matrix. Of course, this makes use of vector multiplication and addition.

F. `asmweights`

The `asmweights` function uses vector operations to perform matrix multiplication between the training images and the eigenvectors, similar to previous functions that use matrix multiplication. The output is stored in a general weight matrix.

G. `eigPart1` and `eigPart2`

Both `eig` functions use matrix multiplication like before through vector operations, with a difference in inputs and outputs. For part 1, the inputs used are the general weight matrix and the image vector of a specific image are used to get a weight vector that will “represent” the specific image. Part 2 uses this weight vector and performs matrix multiplication with the general weight matrix to produce the final eigenface of the image, a matrix with equal size to a normal image vector but weighted based on the general weight matrix

These functions are used not just in the training phase but also in the recognition phase, which occurs right after the training phase ends.

III. Results and Discussion

To compare the execution times (for the training phase) of the C++ and AVX2 implementations of the Eigenface algorithm, the `#define setting` line in the `Eigenface.cpp` file must be set to `true` for the AVX2 implementation and `false` for the C++ implementation. Furthermore, since the execution time of the Eigenface algorithm is dependent on the number of training images, the execution time of both implementations can also be compared with a different number of training images. This can be done by setting the value of `#define trsize` to the desired number of training images.

# of Training Images	Sequential ExecutionTime (ms)	Parallel Execution Time (ms)
4	3272	1340
8	16320	2495
12	65014	18257
16	174213	44077

Table 1.1. Execution Time Comparison Table




Figure 2.1. Execution Time Comparison Graph

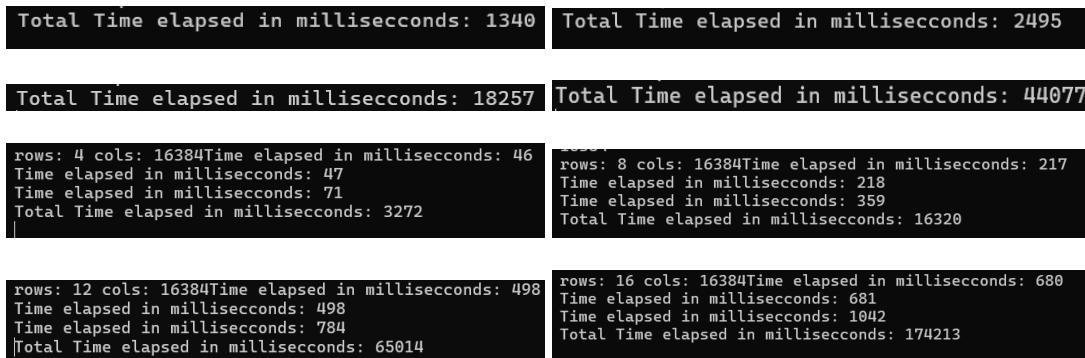


Figure 2.2. Execution Time Captures

From what can be observed from the table and the line graph, the execution time for the C++ implementation is both much higher than the execution time for the AVX2 implementation and grows at a much faster rate. This shows that implementing parallelism to the Eigenface drastically speeds up the program setup.

The accuracy of the recognition face is also affected when using the base C++ algorithm or the AVX2 implementation. Below is a table comparing the relative accuracy of both implementations. The tests are conducted by slowing down the frame capture function of OpenCV and changing the training images to both the tester and no tester, in the same environment, and tallying the “recognition” out of 30 images.

Number of Training Images	C++ Accuracy (out of 30 capture images)	AVX2 Accuracy (out of 30 capture images)
4	24	20
8	21	19
12	23	18
Average Accuracy %	75.556%	63.333%

References

- [1] M.üge Çarıkçı, Figen Özen, *"A Face Recognition System Based on Eigenfaces Method"*, Procedia Technology, Volume 1, 2012.
- [2] M. R. Kawale, Y. Bhadke and V. Inamdar, *"Parallel implementation of eigenface on CUDA,"* 2014 International Conference on Advances in Engineering & Technology Research (ICAE'TR - 2014).
- [3] S. Singh, A. Tripathi, A. Mahajan, and S. Prabhakaran, *"Analysis of Facial Recognition in MATLAB"*. International Journal of Scientific & Engineering Research, February 2012.
- [4] [1 Eigenvalues and eigenvectors](#)